

- Inputs: array of numbers  
e.g.: `poorhousing`  
(1000 houses)
- challenges:
  - class point
  - dimension
  - deformation
  - rotation
  - background changes
  - irrelevant variation
- algorithm:
 

```
def desc(image):
    # magic
    return classifier
```
- solution proposed:
  - edge?
  - convex?
- data-driven
 

```
def extract(img, labels):
    # feature matrix
    # feature model
    def predict(img):
        return class
```
- **first classification: Nearest Neighbors**
  - pixels are features
  - class & labels
  - predict should be able to use these similarly to existing image
  - e.g.: CIFAR10

10 classes  
50,000 training images  
10,000 testing images  
- L1 distance  
 $\| \frac{1}{2} \frac{1}{2} - \frac{1}{1} \|_1 = \left\| \frac{1}{2} \frac{1}{2} \right\|_1 = 15$

```
class NearestNeighbor:
    def __init__(self):
        pass

    def train(self, X, y):
        """ X is N x D where each row is an example. Y is 1-dimension of size N
        A simple nearest neighbor classifier simply remembers all the training data
        self.Xtr = X
        self.ytr = y

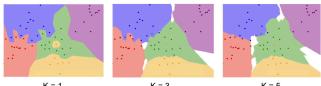
    def predict(self, X):
        """ X is N x D where each row is an example we wish to predict label for """
        num_test = X.shape[0]
        # loop over all test rows
        for i in xrange(num_test):
            # find distances between test row and all training rows
            # using the L1 distance (sum of absolute value differences)
            dists = np.sum(np.abs(Xtr - Xt[i]), axis=1)
            min_index = np.argmin(dists) # index of the minimum distance
            Ypred = self.ytr[min_index] # predict the label of the nearest neighbor
        return Ypred
```

• **@ k NN**  
hard cases should be hard @ prediction class @ learning



#### - K-NN

K-Nearest Neighbors  
Instead of copying label from nearest neighbor,  
take majority vote from K closest points



- difference metric  
 $L_1 = \text{manhattan}$ ,  $L_2 = \text{euclidean}$
- hyperparameter
  - the parameters are static training
  - split data into  $\begin{cases} \text{train} \\ \text{val} \\ \text{test} \end{cases}$  choose hyperparameters on val  $\forall$   $\begin{cases} \text{train} \\ \text{val} \\ \text{test} \end{cases}$
  - cross-validation

#### Setting Hyperparameters

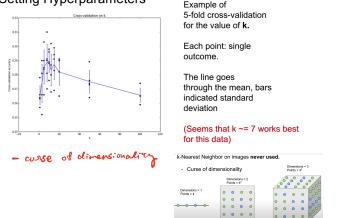
Your Dataset

Idea #4: Cross-Validation: Split data into folds,  
try each fold as validation and average the results

fold 1	fold 2	fold 3	fold 4	fold 5	test
fold 1	fold 2	fold 3	fold 4	fold 5	test
fold 1	fold 2	fold 3	fold 4	fold 5	test

Useful for small datasets, but not used too frequently in deep learning

#### Setting Hyperparameters



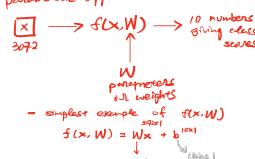
#### • Linear classifiers

##### - Neural Network

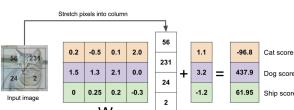


- CIFAR10  
50,000 train (32x32x3)  
10,000 test 3072

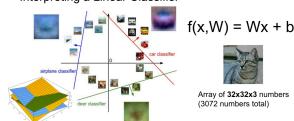
##### - parametric approach



Example with an image with 4 pixels, and 3 classes (cat/dog/ship)



#### Interpreting a Linear Classifier



#### Hard cases for a linear classifier

Class 1: number of pixels > 0 odd

Class 1: number of pixels > 0 even

Class 1:  $\|x\|_2 \leq 2$

Class 1:  $\|x\|_2 \geq 2$

Class 1: everything else

Class 1: everything else

Class 1: three modes

Class 1: everything else

Class 1: everything else