

- inputs: array of numbers
e.g. $100 \times 100 \times 3$ (RGB channels)
- challenges:
 - view point
 - illumination
 - occlusion
 - background clutter
 - intrinsics variation
- algorithm:


```
def classifyImage(image):
    # logic
    return classifier
```
- solution proposal:
 - edge?
 - corner?
- data-driven


```
def training(image, labels):
    # Machine learning!
    return model
def predict(model, image):
    return class
```
- gross classifier: Neural Network
 - parametric fit
 - done & stable
 - predict the label of the most similar training image
- e.g. CIFAR-10
 - 10 classes
 - 50,000 training imgs
 - 10,000 testing imgs
- L1 distance
 $| \frac{1}{2} \frac{3}{4} - 1 \cdot | = | \frac{1}{2} \frac{3}{4} |_1 = 1.5$

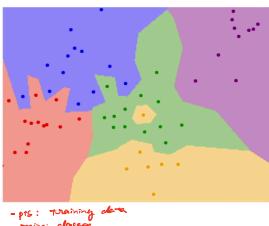
```
class NearestNeighbor:
    def __init__(self):
        pass

    def train(self, X, y):
        """ X is N x D where each row is an example. Y is 1-dimension of size N """
        # the nearest neighbor classifier simply remembers all the training data
        self.Xtr = X
        self.ytr = y

    def predict(self, X):
        """ X is K x D where each row is an example we wish to predict label for """
        num_test = X.shape[0]
        # lets make sure that the output type matches the input type
        Ypred = np.zeros(num_test, dtype = self.ytr.dtype)

        for i in range(num_test):
            if i % 100 == 0:
                print('predicting %d' % i)
            # find the nearest training image to the i-th test image
            # a vector of shape N where each row is a difference in value differences
            distances = np.sum(np.abs(self.Xtr - X[i, :]), axis = 1)
            min_index = np.argmin(distances) # get the index with smallest distance
            Ypred[i] = self.ytr[min_index] # predict the label of the nearest example
        return Ypred
```

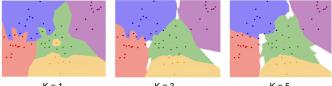
- @ w/ N examples
how fast is it?
At Train O(1)
predict O(N)
(bad classifier should be fast @ prediction slow @ training)



K-NN

K-Nearest Neighbors

Instead of copying label from nearest neighbor,
take majority vote from K closest points



- hyperparameters
 - the parameters are their training
 - split data into {choose hyperparameters on validation set}
 - train
 - test
- cross-validation

Setting Hyperparameters

Your Dataset				
fold 1	fold 2	fold 3	fold 4	fold 5
fold 1	fold 2	fold 3	fold 4	fold 5
fold 1	fold 2	fold 3	fold 4	fold 5

Idea #4: Cross-Validation: Split data into folds, try each fold as validation and average the results

fold 1	fold 2	fold 3	fold 4	fold 5	test
fold 1	fold 2	fold 3	fold 4	fold 5	test
fold 1	fold 2	fold 3	fold 4	fold 5	test

Useful for small datasets, but not used too frequently in deep learning.

Setting Hyperparameters

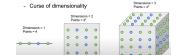
Example of 5-fold cross-validation for the value of k.

Each point: single outcome.

The line goes through the mean, bars indicated standard deviation

(Seems that $k \approx 7$ works best for this data)

- curse of dimensionality



Linear classifiers

Neural Network



CIFAR-10

50,000 train (32x32x3)
10,000 test 3072

parametric approach

$$\boxed{x} \rightarrow f(x, W) \rightarrow \begin{matrix} 10 \text{ numbers} \\ \text{giving class scores} \end{matrix}$$

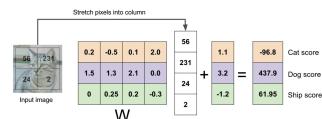
W parameter
6x3 weights

simplest example of $f(x, W)$

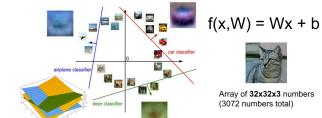
$$f(x, W) = Wx + b^{(10)}$$

$10 \times 3 \times 3$

Example with an image with 4 pixels, and 3 classes (cat/dog(ship))



Interpreting a Linear Classifier



Hard cases for a linear classifier

Class 1:
number of pixels > 0 odd

Class 2:
number of pixels > 0 even

Class 1:
 $\|x\|_2 \leq 2$

Class 2:
Everything else

point problem (true/false)

multi-modal

centered

