

COMP5212: Machine Learning

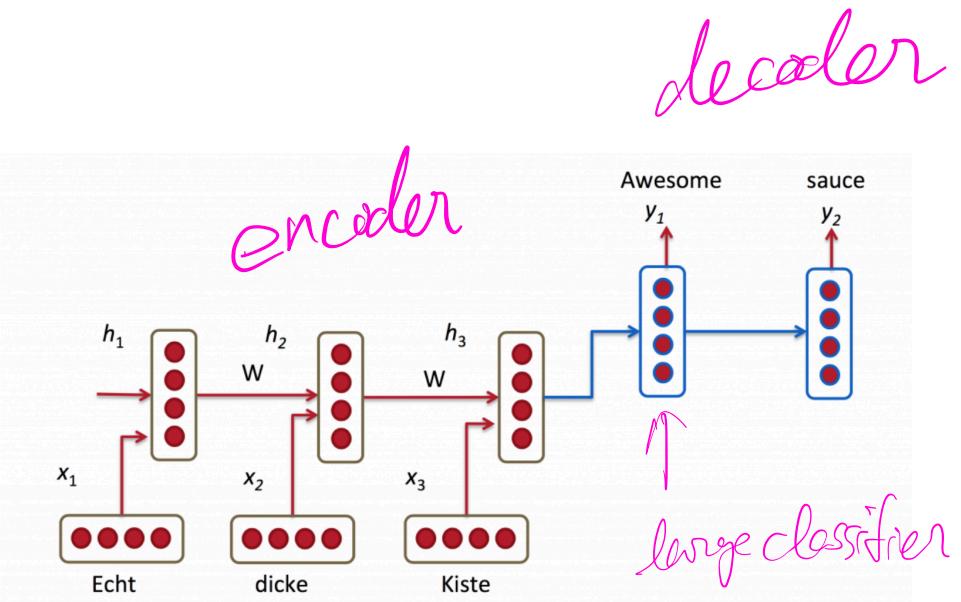
Lecture 15

Minhao Cheng

Recurrent Neural Network

Neural Machine Translation (NMT)

- Output the translated sentence from an input sentence
- Training data: a set of input-output pairs (supervised setting)
- Encoder-decoder approach:
 - Encoder: Use (RNN/LSTM) to encode the input sentence into a latent vector
 - Decoder: Use (RNN/LSTM) to generate a sentence based on the latent vector

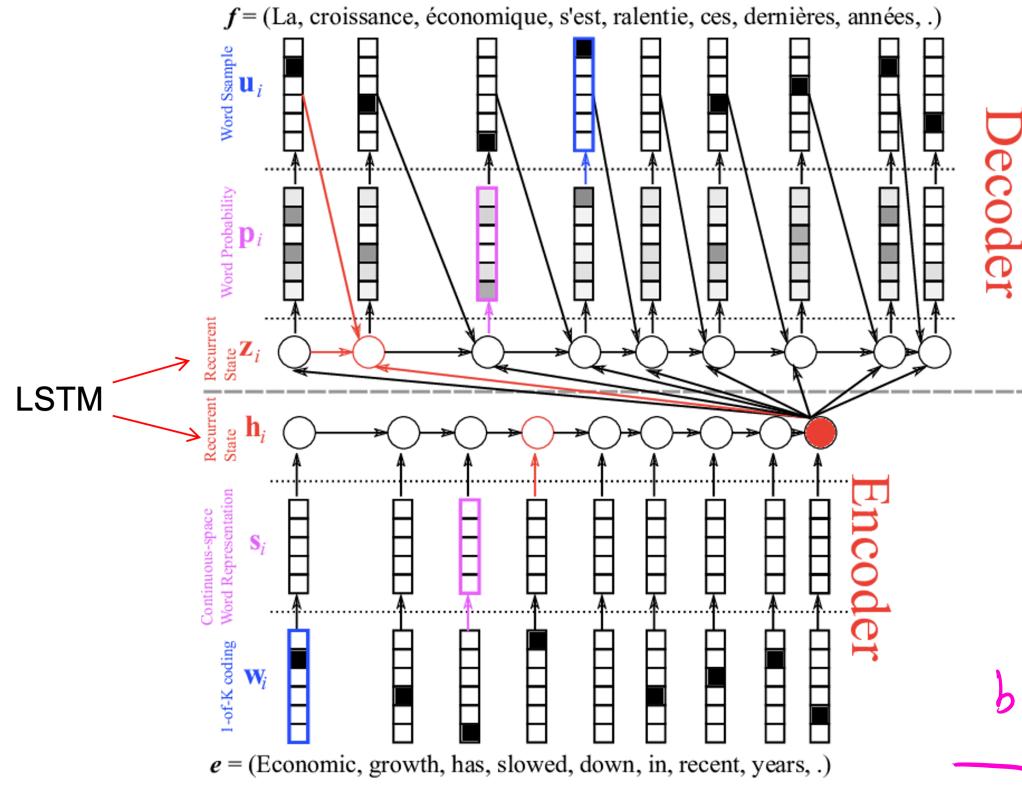


CNN for NLP

word \rightarrow embedding
of words

Recurrent Neural Network

Neural Machine Translation



attention in
NLP

bi-LSTM

↳ which words hv higher connectivity, that we could use

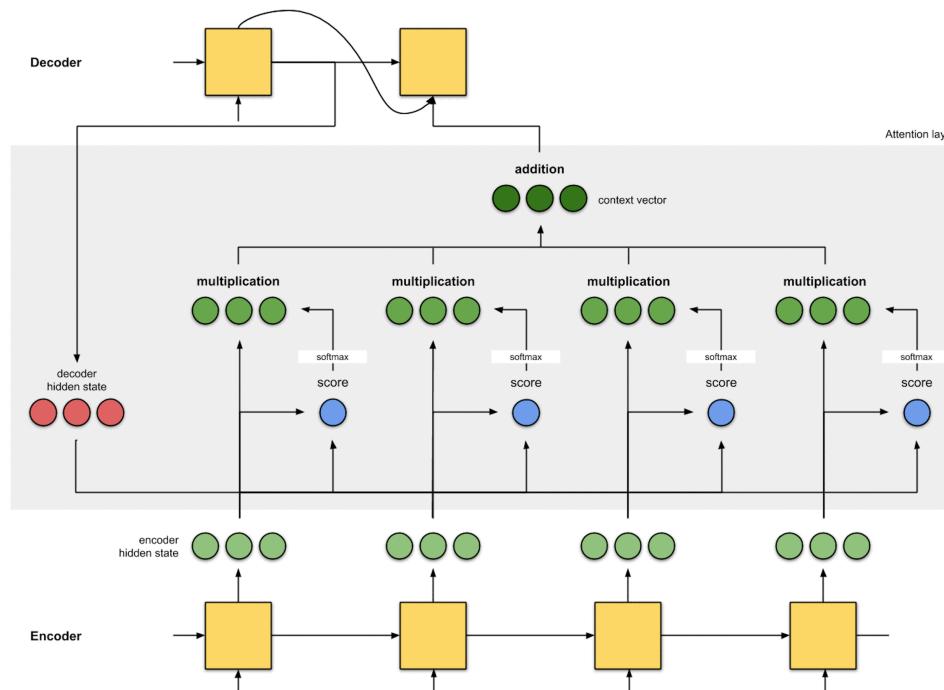
Recurrent Neural Network

Attention in NMT

- Usually, each output word is only related to a subset of input words (e.g., for machine translation)
- Let u be the current decoder latent state, v_1, \dots, v_n be the latent state for each input word
- Compute the weight of each state by
 - $p = \text{Softmax}(u^T v_1, \dots, u^T v_n)$ calculate the attention
- Compute the context vector by $Vp = p_1 v_1 + \dots + p_n v_n$

Recurrent Neural Network

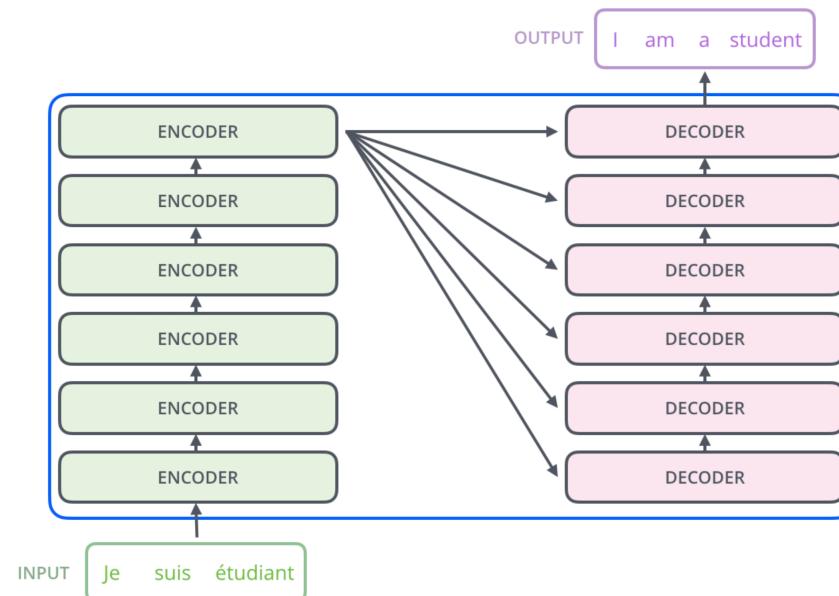
Attention in NMT



Transformer

Transformer

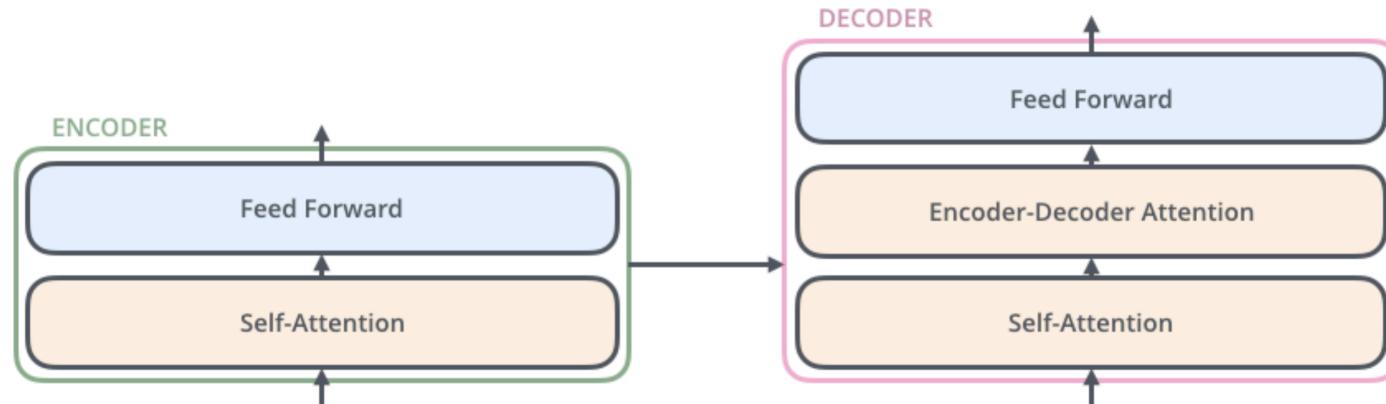
- An architecture that replies entirely on attention without using CNN/RNN
- Proposed in "Attention Is All You Need" (Vaswani et al., 2017)
- Initially used for neural machine translation



Transformer

Encoder and Decoder

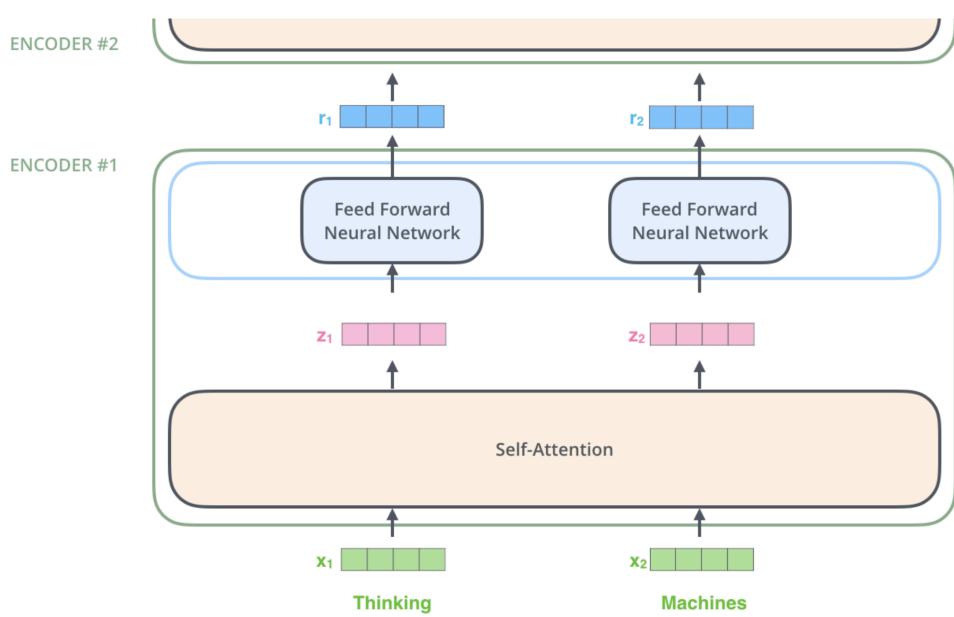
- Self attention layer: the main architecture used in Transformer
- Decoder: will have another attention layer to help it focuses on relevant parts of input sentences.



Transformer

Encoder

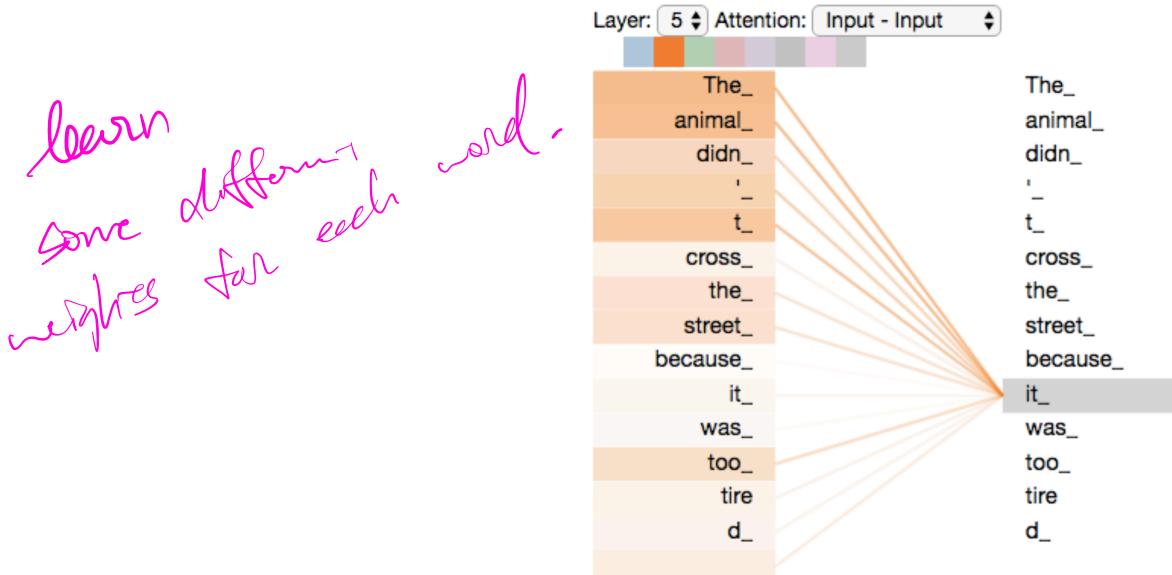
- Each word has a corresponding "latent vector" (initially the word embedding for each word)
- Each layer of encoder:
 - Receive a list of vectors as input
 - Passing these vectors to a **self-attention** layer
 - Then passing them into a feed-forward layer
 - Output a list of vectors



Transformer

Self-attention layer

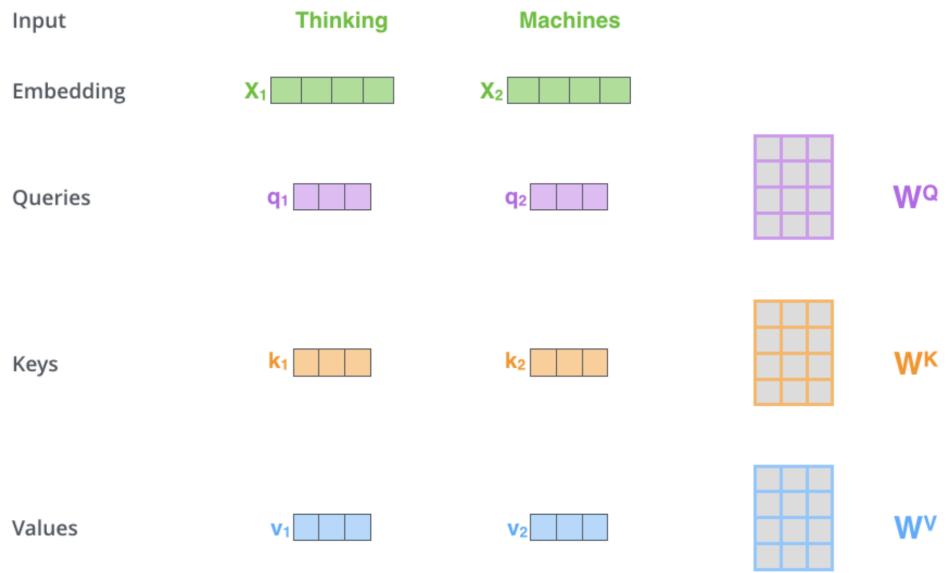
- Main idea: The actual meaning of each word may be related to other words in the sentence
- The actual meaning (latent vector) of each word is a weighted (attention) combination of other words (latent vectors) in the sentences



Transformer

Self-attention layer

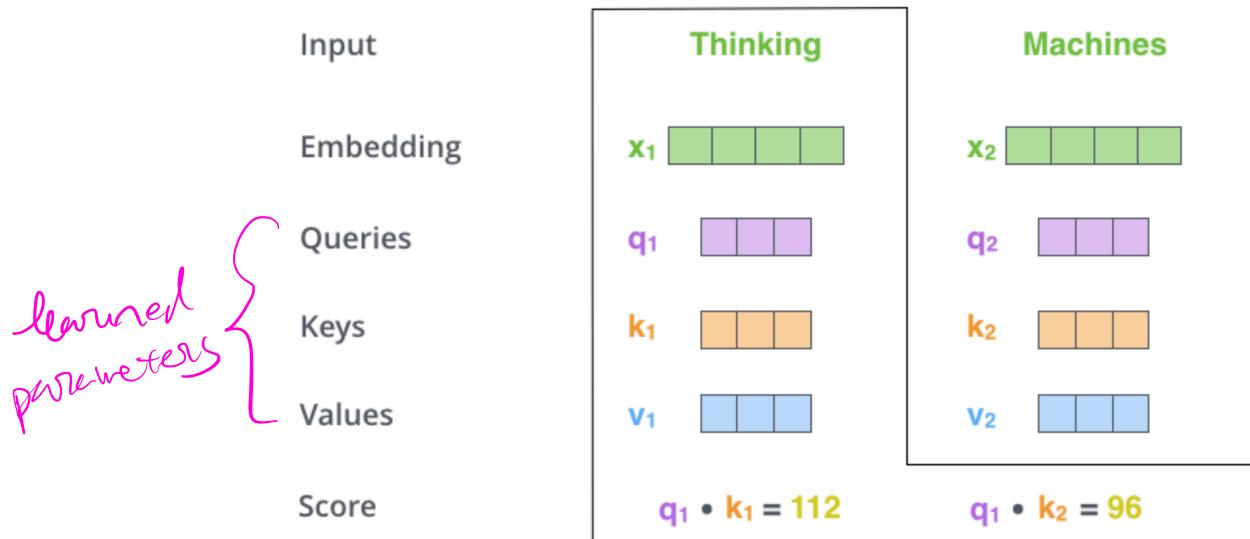
- Input latent vectors: x_1, \dots, x_n
- Self-attention parameters:
 W^Q, W^K, W^V (weights for query, key, value)
- For each word i , compute
 - Query vector: $q_i = x_i W^Q$
 - Key vector: $k_i = x_i W^K$
 - Value vector: $v_i = x_i W^V$



Transformer

Self-attention layer

- For each word i , compute the scores to determine how much focus to place on other input words
 - The attention score for word j to word i : $q_i^T k_j$



Transformer

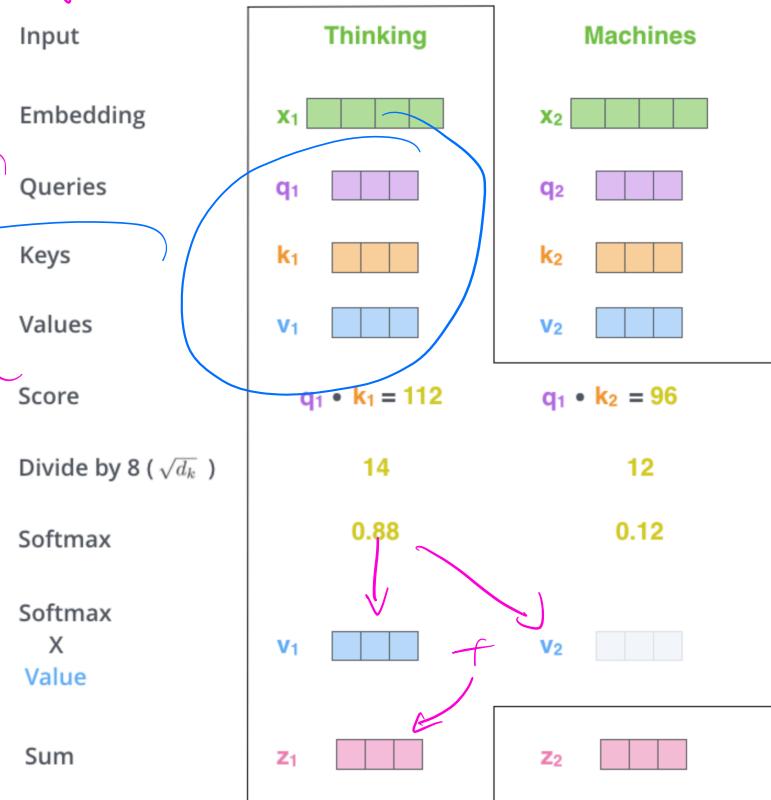
Self-attention layer

using padding to deal with different length, so that the input could be the same size

- For each word i , the output vector

- $\sum_j s_{ij} v_j, \quad s_i = \text{softmax}(q_i^T k_1, \dots, q_i^T k_n)$

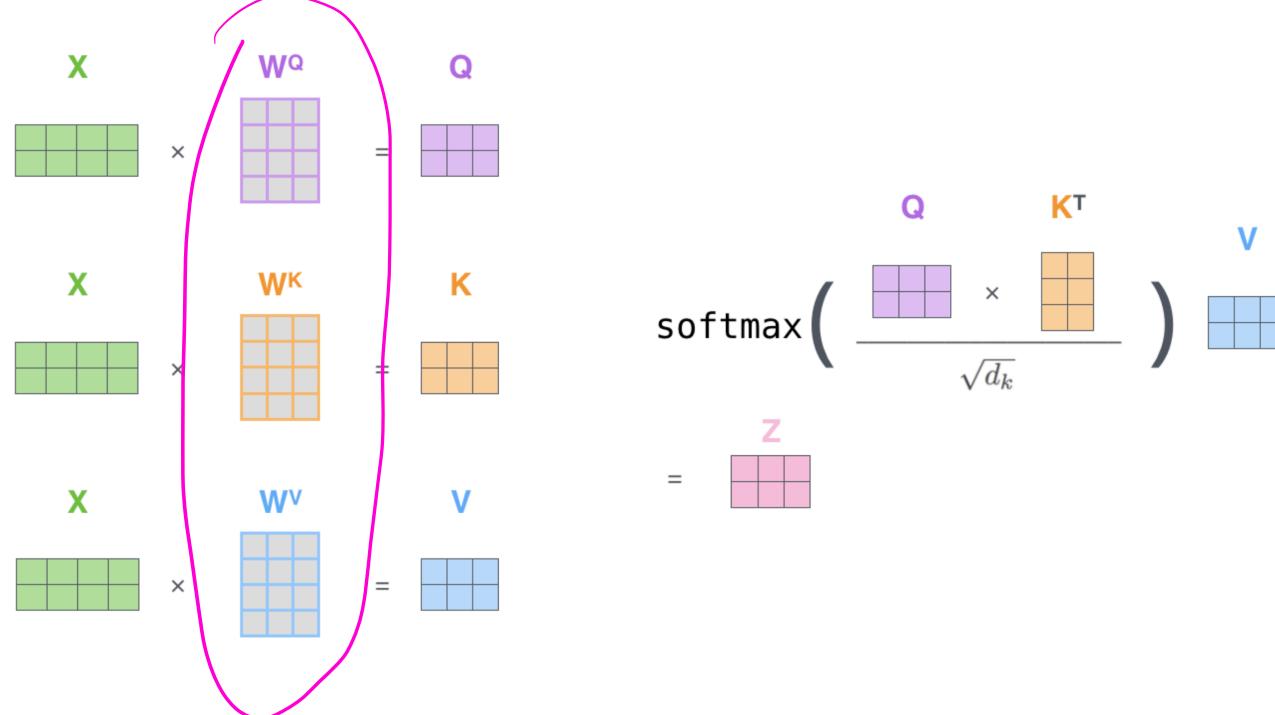
could
be
 8×3 , 16×3
into contains info



Transformer

Matrix form

- $Q = XW^Q, K = XW^K, V = XW^V, Z = \text{softmax}(QK^T)V$



Transformer

Multiply with weight matrix to reshape

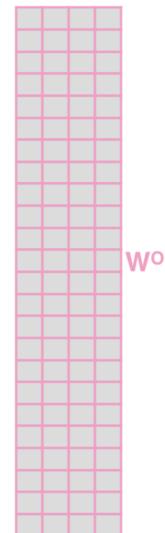
- Gather all the outputs Z_1, \dots, Z_k
- Multiply with a weight matrix to reshape
- Then pass to the next fully connected layer

1) Concatenate all the attention heads



2) Multiply with a weight matrix W^o that was trained jointly with the model

x



3) The result would be the Z matrix that captures information from all the attention heads. We can send this forward to the FFNN

$$= \begin{matrix} & Z \\ & \begin{matrix} & & \\ & & \\ & & \end{matrix} \end{matrix}$$

Transformer

Overall architecture

- 1) This is our input sentence*
2) We embed each word*

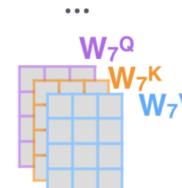
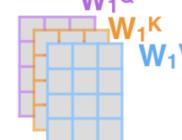
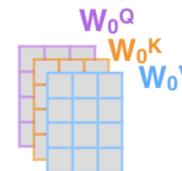
Thinking Machines
 X

* In all encoders other than #0, we don't need embedding. We start directly with the output of the encoder right below this one

R

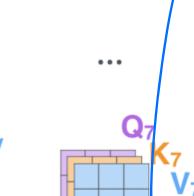
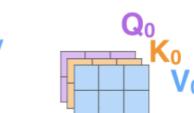
Self - attention

- 3) Split into 8 heads. We multiply X or R with weight matrices

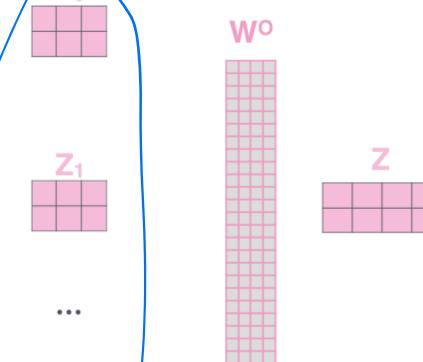


8 heads
attention

- 4) Calculate attention using the resulting $Q/K/V$ matrices



- 5) Concatenate the resulting Z matrices, then multiply with weight matrix W^o to produce the output of the layer



fuse
all info
into
smaller |

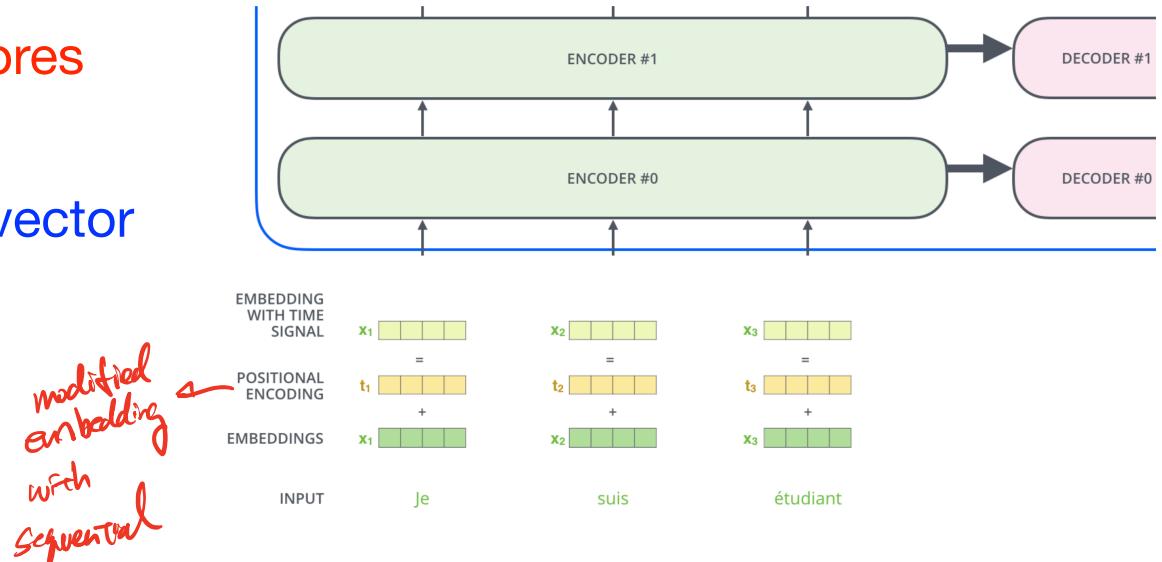
can be trained
much faster than RNN,
as a lot of weights are
in parallel

Transformer

Sinusoidal Position Encoding

- The above architecture **ignores the sequential information**
- Add a **positional encoding vector** to each x_i (according to i)

→ to enforce sequential info



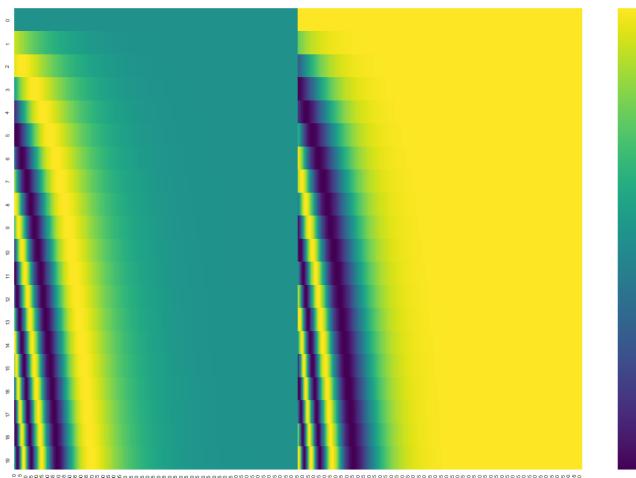
Transformer

Positional Embedding

- Sin/cosine functions with different wavelengths (used in the original Transformer)

- The jth dimension of ith token $p_i[j] = \begin{cases} \sin(i \cdot c^{\frac{j}{d}}) & \text{if } j \text{ is even} \\ \cos(i \cdot c^{\frac{j-1}{d}}) & \text{if } j \text{ is odd} \end{cases}$

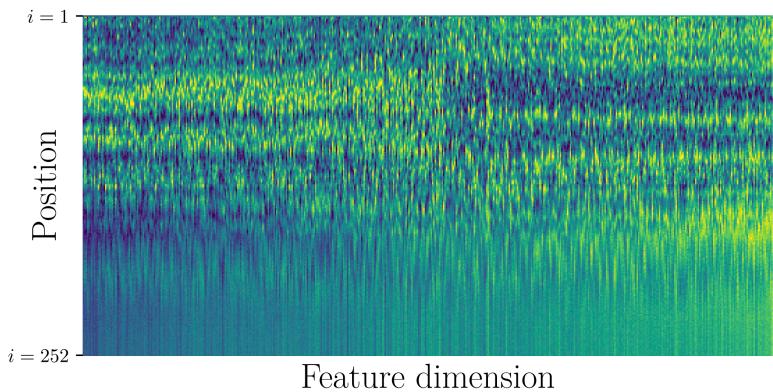
- smooth, parameter-free, inductive



Transformer

Types of positional encoding

- Position embedding: learn a latent vector for each position
 - non-smooth, data-driven (learnable), non-inductive
- Relative position embedding:
 - For each i, j , use the relative position embedding a_{j-i}
 - non-smooth, data-driven (learnable), (partial)-inductive



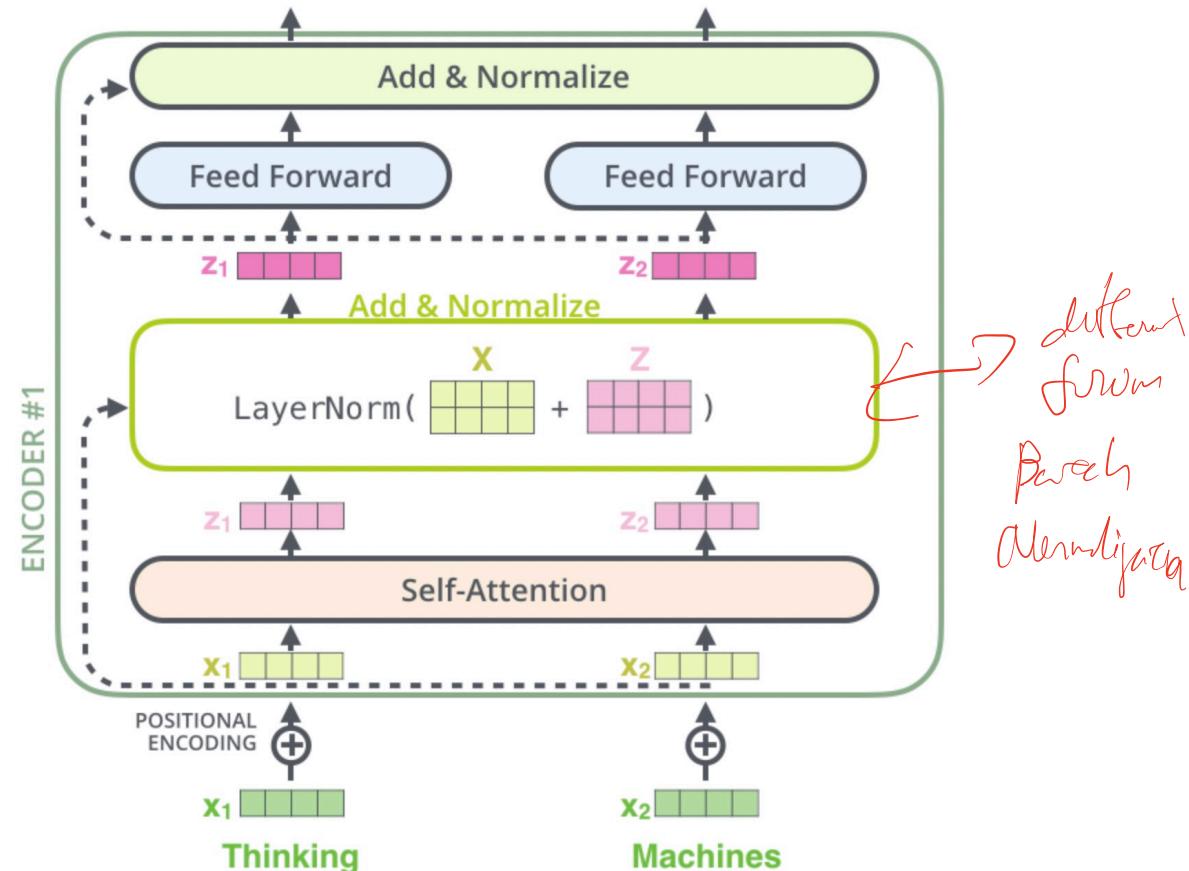
Transformer

Positional Encoding

- Neural ODE embedding :
 - Model positional embedding as a dynamic linear system
 - “Learning to Encode Position for Transformer with Continuous Dynamical Model. Liu et al., 2020”
- Learnable Fourier Feature (Li et al., 2021):
 - $p_x = \phi(r_x, \theta)W_p$, where $r_x = \frac{1}{D}[\cos xW_r, \sin xW_r]$
 - W_r, W_p : learnable parameters (irrelevant to sequence length)
 - “Learnable Fourier Features for Multi-Dimensional Spatial Positional Encoding. Li et al., 2021”
 - smooth, **data-driven (learnable)**, inductive

Transformer

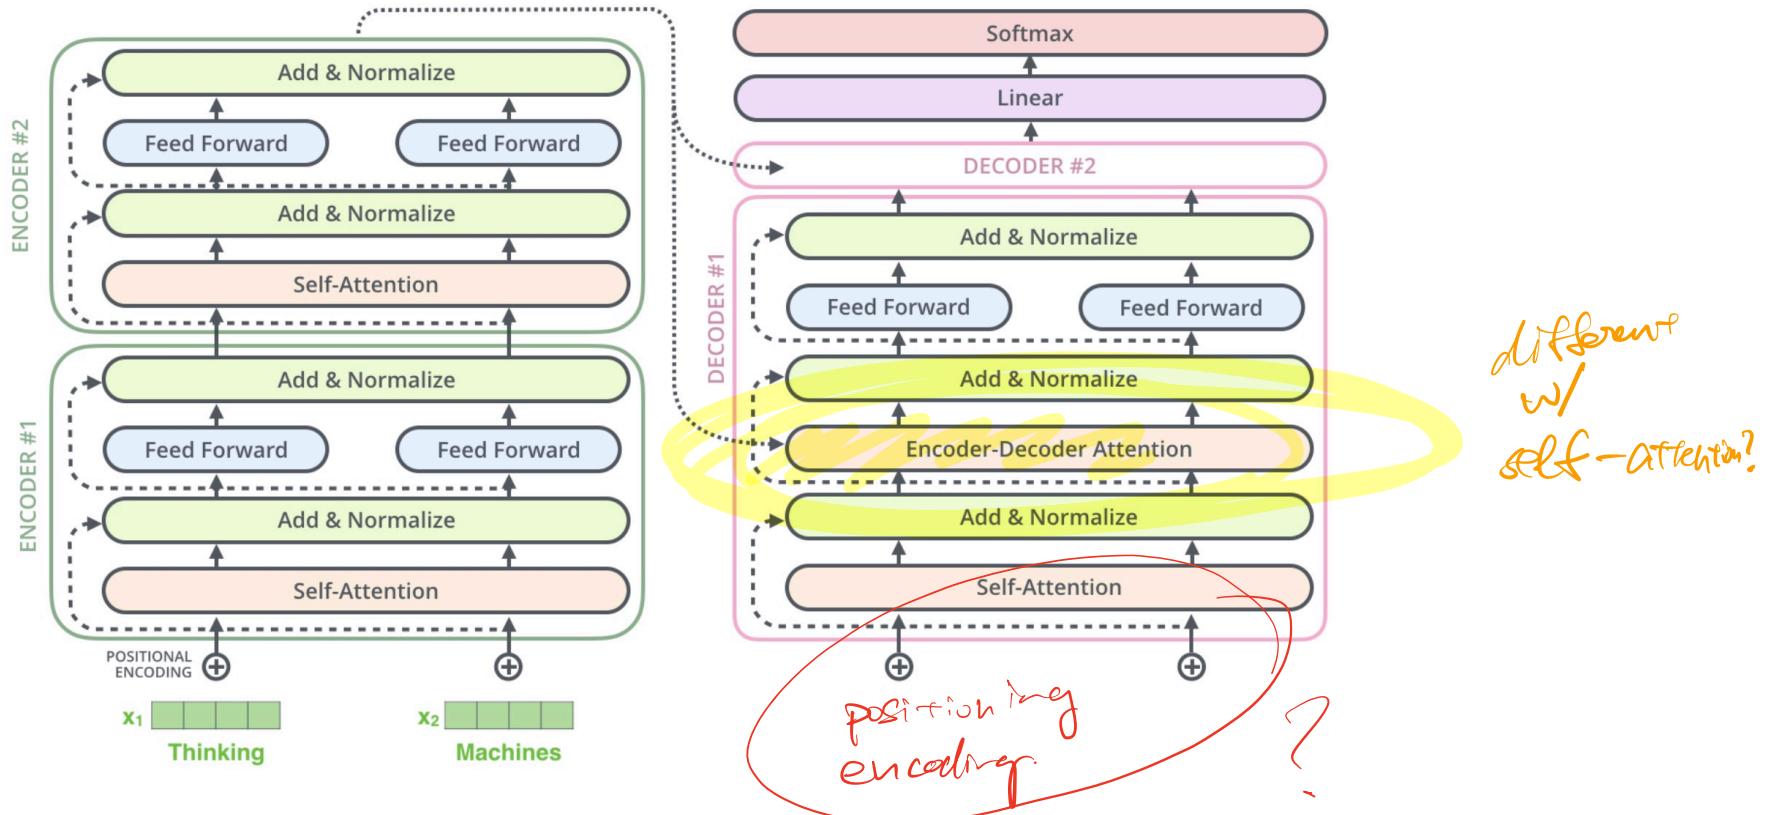
Residual



Transformer

Whole framework

Just matrix multiplication!
very simple



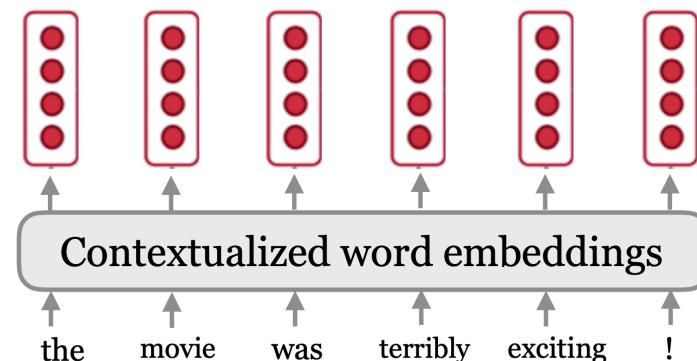
Contextual embedding

Contextual word representation

- The semantic meaning of a word should depend on its context



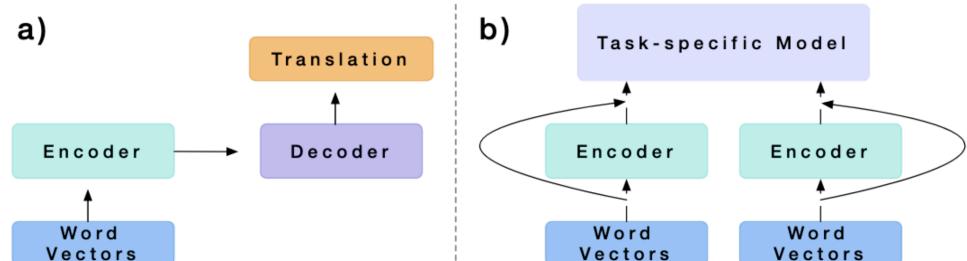
- Solution: Train a model to extract contextual representations on text corpus



Contextual embedding

CoVe (McCann et al., 2017)

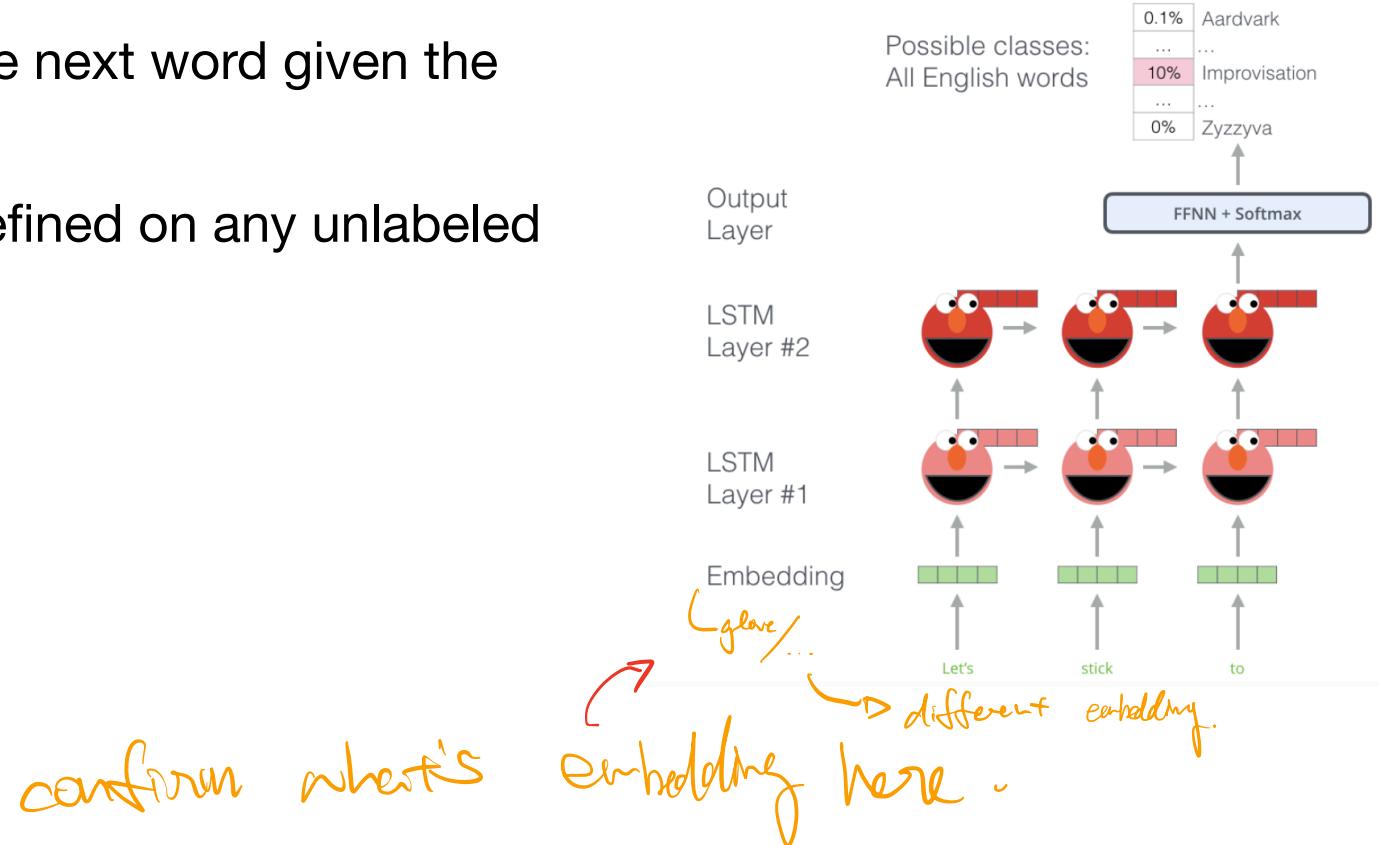
- Key idea: Train a standard neural machine translation model
- Take the encoder directly as contextualized word embeddings
- Problems:
 - Translation requires paired (labeled) data
 - The embeddings are tailored to particular translation corpuses



Contextual embedding

Language model pretraining task

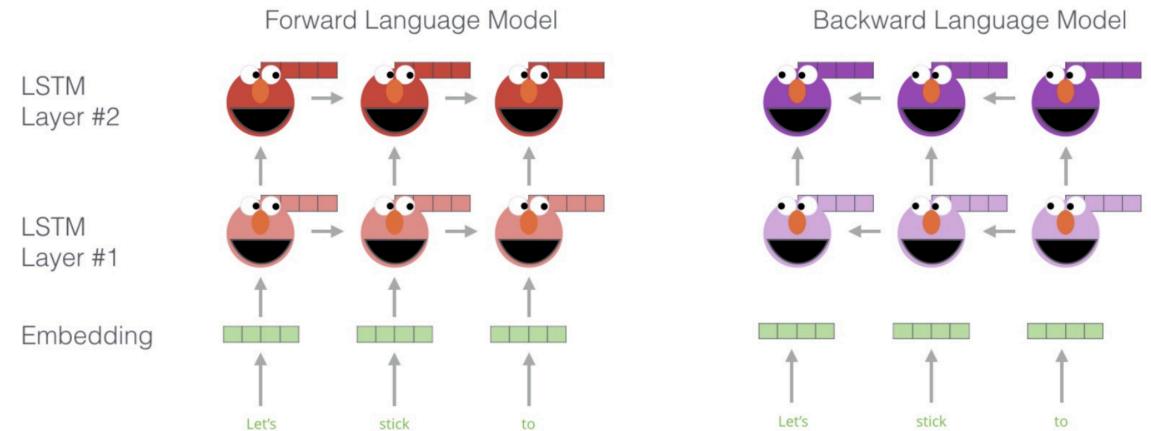
- Predict the next word given the prefix
- Can be defined on any unlabeled document



Contextual embedding

ELMo (Peter et al., 2018)

- Key ideas:
- Train a forward and backward LSTM language model on large corpus
- Use the hidden states for each token to compute a vector representation of each word
- Replace the word embedding by Elmo's embedding (with fixed Elmo's LSTM weights)



Contextual embedding

ELMo results

Task	Previous SOTA		Our Baseline	ELMo + Baseline	Increase (Absolute/Relative)
SQuAD	Liu et al. (2017)	84.4	81.1	85.8	4.7 / 24.9%
SNLI	Chen et al. (2017)	88.6	88.0	88.7 ± 0.17	0.7 / 5.8%
SRL	He et al. (2017)	81.7	81.4	84.6	3.2 / 17.2%
Coref	Lee et al. (2017)	67.2	67.2	70.4	3.2 / 9.8%
NER	Peters et al. (2017)	91.93 ± 0.19	90.15	92.22 ± 0.10	2.06 / 21%
SST-5	McCann et al. (2017)	53.7	51.4	54.7 ± 0.5	3.3 / 6.8%

Contextual embedding

BERT

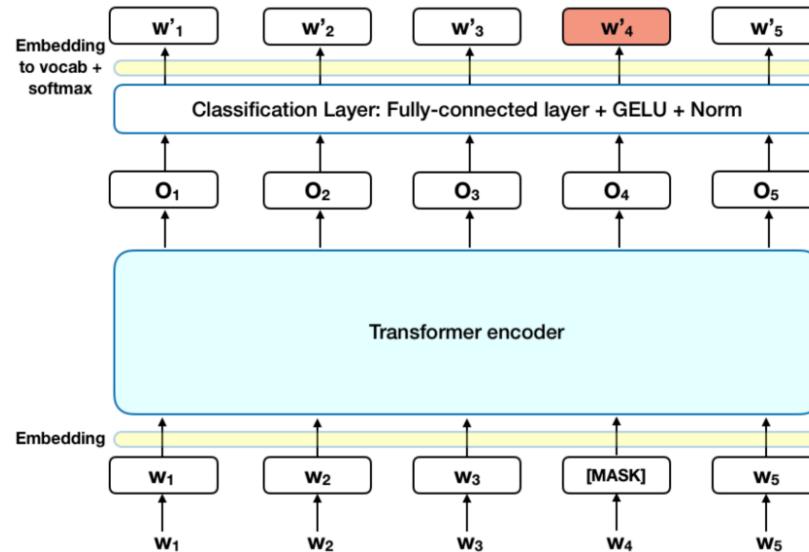
CBOW
Skip Gram

- Key idea: replace LSTM by Transformer
- Define the generated pretraining task by masked language model
- Two pretraining tasks
- Finetune both BERT weights and task-dependent model weights for each task

Contextual embedding

BERT pretraining loss

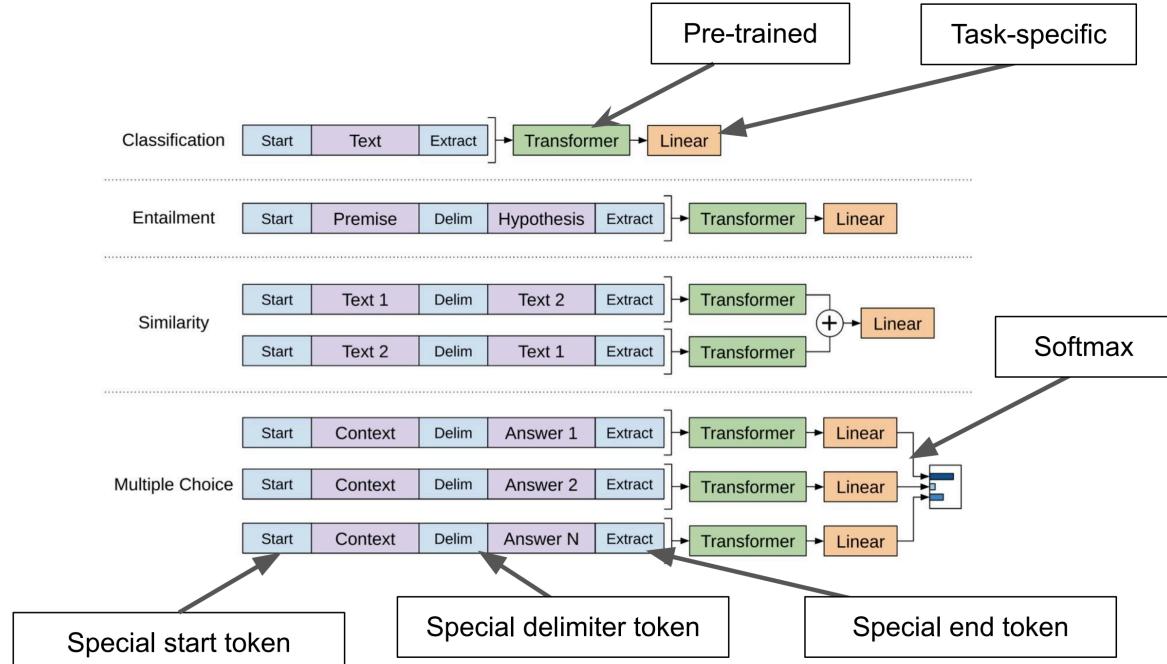
- Masked language model: predicting each word by the rest of sentence
- Next sentence prediction: the model receives pairs of sentences as input and learns to predict if the second sentence is the subsequent sentence in the original document.



Contextual embedding

BERT finetuning

- Keep the pretrained Transformers
- Replace or append a layer for the final task
- Train the whole model based on the task-dependent loss



Contextual embedding

BERT results

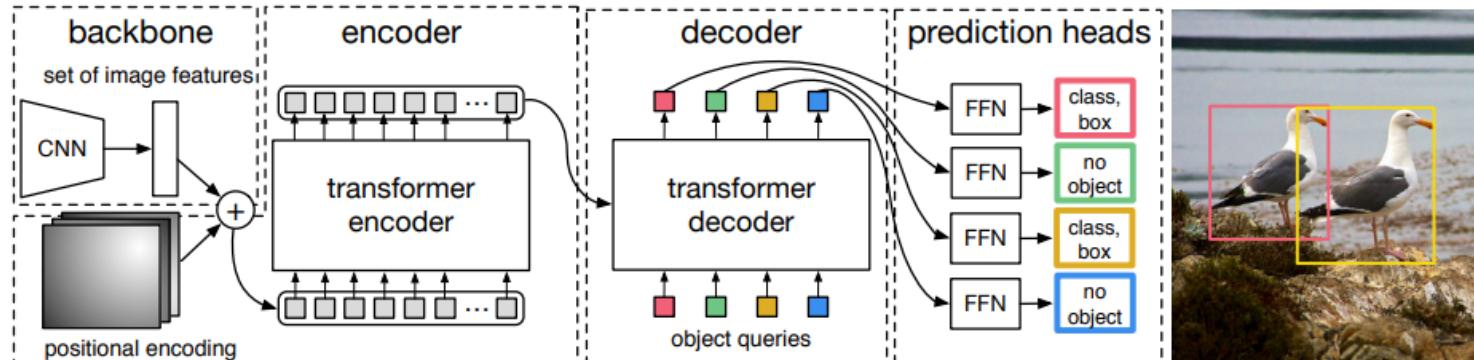
System	MNLI-(m/mm) 392k	QQP 363k	QNLI 108k	SST-2 67k	CoLA 8.5k	STS-B 5.7k	MRPC 3.5k	RTE 2.5k	Average
Pre-OpenAI SOTA	80.6/80.1	66.1	82.3	93.2	35.0	81.0	86.0	61.7	74.0
BiLSTM+ELMo+Attn	76.4/76.1	64.8	79.8	90.4	36.0	73.3	84.9	56.8	71.0
OpenAI GPT	82.1/81.4	70.3	87.4	91.3	45.4	80.0	82.3	56.0	75.1
BERT _{BASE}	84.6/83.4	71.2	90.5	93.5	52.1	85.8	88.9	66.4	79.6
BERT _{LARGE}	86.7/85.9	72.1	92.7	94.9	60.5	86.5	89.3	70.1	82.1

What is the metric here ?

Vision Transformer (ViT)

Attempts on applying self-attention to vision

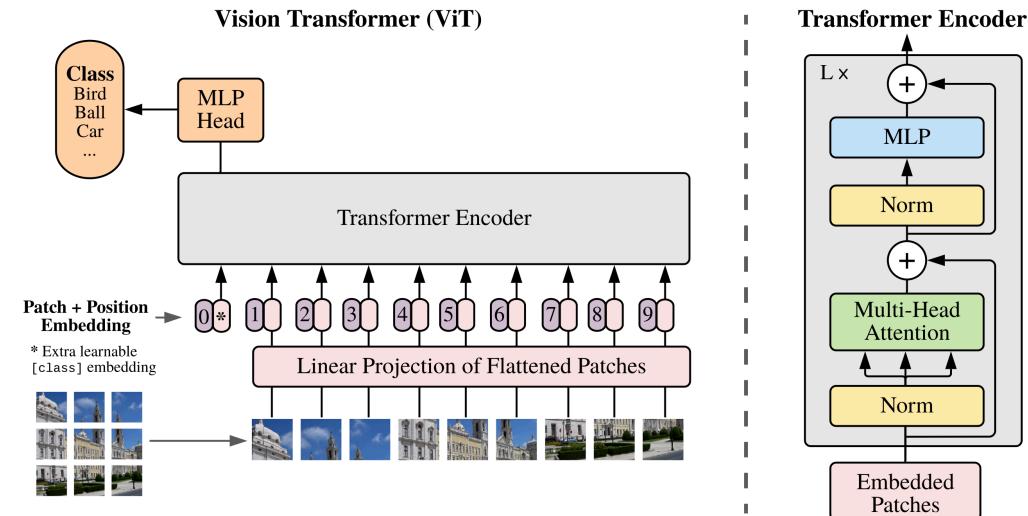
- DETR (Carion et al., 2020): CNN + Self-attention for object detection
- Stand-alone self-attention (Ramachandran et al., 2020)
- ...



Vision Transformer (ViT)

Vision Transformer (ViT)

- Partition input image into $K \times K$ patches
- A linear projection to transform each patch to feature (no convolution)
- Pass tokens into Transformer



Vision Transformer (ViT)

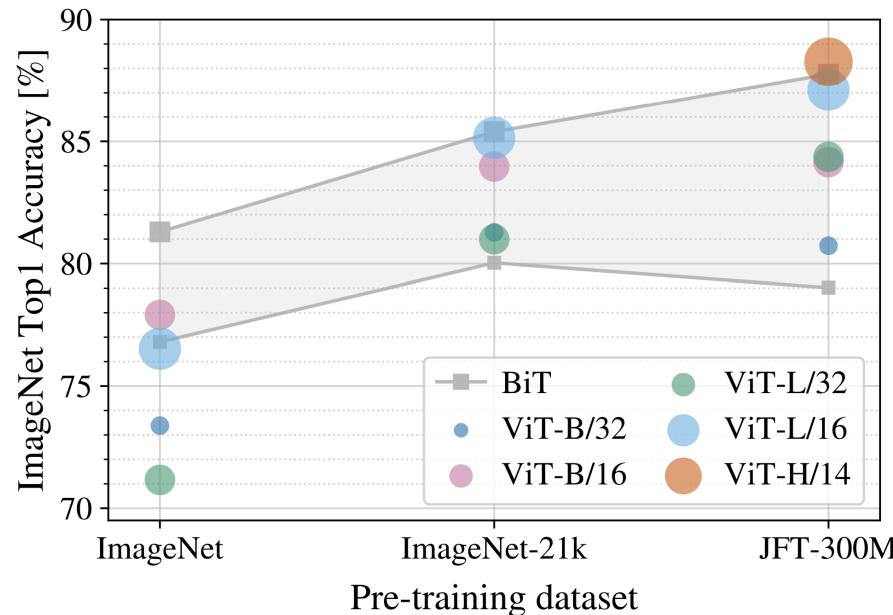
Vision Transformer (ViT)

- Patches are non-overlapping in the original ViT
- $N \times N$ image $\Rightarrow (N/K)^2$ tokens
- Smaller patch size \Rightarrow more input tokens
 - Higher computation (memory) cost, (usually) higher accuracy
- Use 1D (learnable) positional embedding
- Inference with higher resolution:
 - Keep the same patch size, which leads to longer sequence
 - Interpolation for positional embedding

Vision Transformer (ViT)

ViT Performance

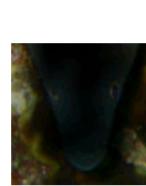
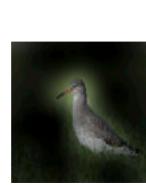
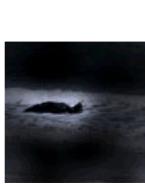
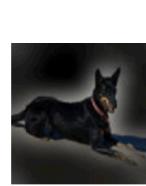
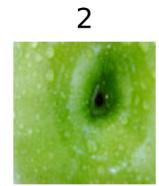
- ViT outperforms CNN with large pretraining



Vision Transformer (ViT)

ViT Performance

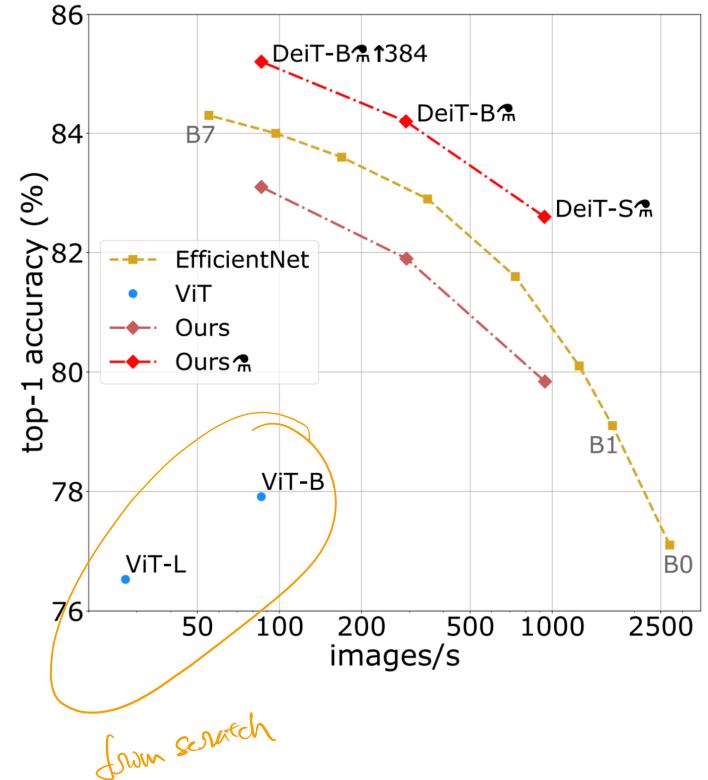
- Attention maps of ViT (to input)



Vision Transformer (ViT)

ViT v.s. ResNet

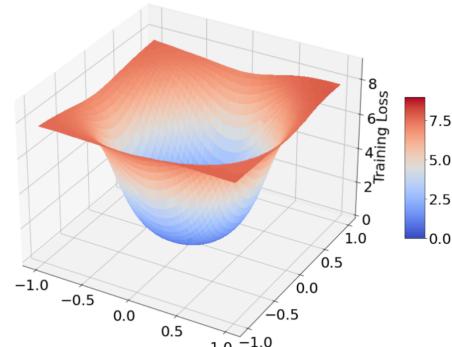
- Can ViT outperform ResNet on ImageNet without pretraining?
- DeiT (Touvron et al., 2021):
 - Use very strong data augmentation
 - Use a ResNet teacher and distill to ViT



Vision Transformer (ViT)

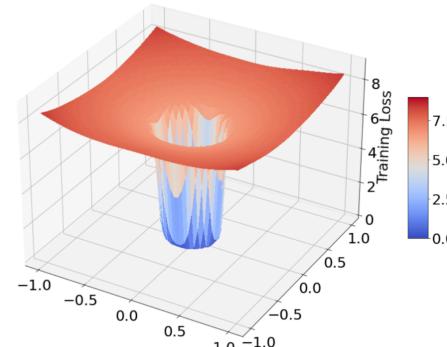
ViT v.s. ResNet

- ViT tends to converge to sharper regions than ResNet



(a) ResNet

Leading eigenvalue of
Hessian: **179.8**



(b) ViT

Leading eigenvalue of
Hessian: **738.8**

Vision Transformer (ViT)

“Sharpness” is related to generalization

- Testing can be viewed as a slightly perturbed training distribution
- Sharp minimum \Rightarrow performance degrades significantly from training to testing

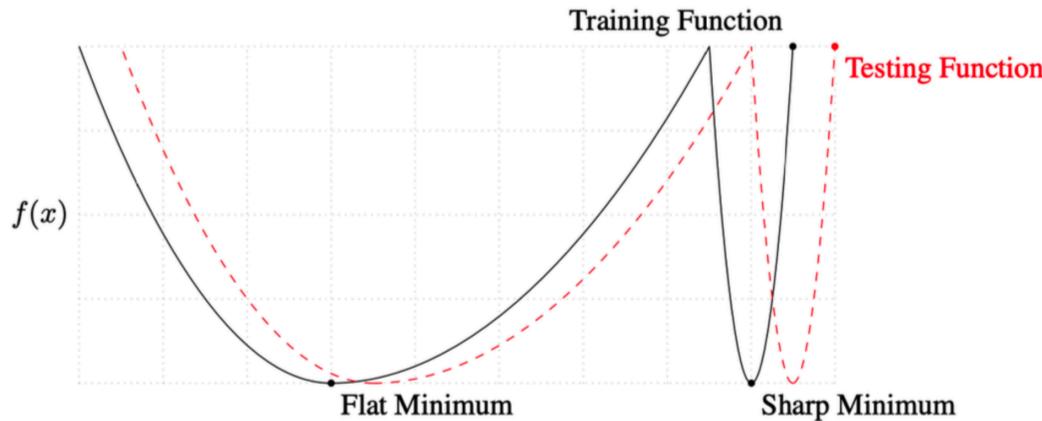


Figure from (Keskar et al., 2017)

Vision Transformer (ViT)

Sharpness Aware Minimization (SAM)

- Optimize the worst-case loss within a small neighborhood

- $\min_w \max_{\|\delta\|_2 \leq \epsilon} L(w + \delta)$

- ϵ is a small constant (hyper-parameter)

- Use 1-step gradient ascent to approximate inner max:

- $\hat{\delta} = \arg \max_{\|\delta\|_2 \leq \epsilon} L(w) + \nabla L(w)^T \delta = \epsilon \frac{\nabla L(w)}{\|\nabla L(w)\|}$

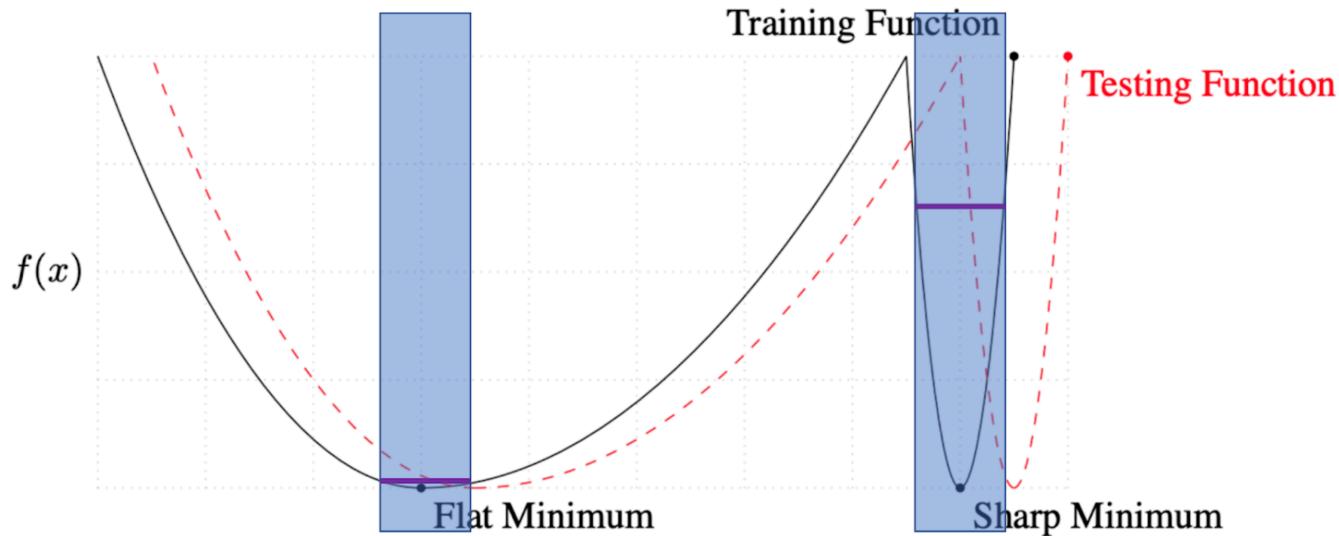
- Conduct the following update for each iteration:

- $w \leftarrow w - \alpha \nabla L(w + \hat{\delta})$

Vision Transformer (ViT)

Sharpness Aware Minimization (SAM)

- SAM is a natural way to penalize sharpness region (but requires some computational overhead)



Vision Transformer (ViT)

ViT v.s. ResNet

Model	#params	Throughput (img/sec/core)	ImageNet	Real	V2	ImageNet-R	ImageNet-C
ResNet							
ResNet-50-SAM	25M	2161	76.7 (+0.7)	83.1 (+0.7)	64.6 (+1.0)	23.3 (+1.1)	46.5 (+1.9)
ResNet-101-SAM	44M	1334	78.6 (+0.8)	84.8 (+0.9)	66.7 (+1.4)	25.9 (+1.5)	51.3 (+2.8)
ResNet-152-SAM	60M	935	79.3 (+0.8)	84.9 (+0.7)	67.3 (+1.0)	25.7 (+0.4)	52.2 (+2.2)
ResNet-50x2-SAM	98M	891	79.6 (+1.5)	85.3 (+1.6)	67.5 (+1.7)	26.0 (+2.9)	50.7 (+3.9)
ResNet-101x2-SAM	173M	519	80.9 (+2.4)	86.4 (+2.4)	69.1 (+2.8)	27.8 (+3.2)	54.0 (+4.7)
ResNet-152x2-SAM	236M	356	81.1 (+1.8)	86.4 (+1.9)	69.6 (+2.3)	28.1 (+2.8)	55.0 (+4.2)
Vision Transformer							
ViT-S/32-SAM	23M	6888	70.5 (+2.1)	77.5 (+2.3)	56.9 (+2.6)	21.4 (+2.4)	46.2 (+2.9)
ViT-S/16-SAM	22M	2043	78.1 (+3.7)	84.1 (+3.7)	65.6 (+3.9)	24.7 (+4.7)	53.0 (+6.5)
ViT-S/14-SAM	22M	1234	78.8 (+4.0)	84.8 (+4.5)	67.2 (+5.2)	24.4 (+4.7)	54.2 (+7.0)
ViT-S/8-SAM	22M	333	81.3 (+5.3)	86.7 (+5.5)	70.4 (+6.2)	25.3 (+6.1)	55.6 (+8.5)
ViT-B/32-SAM	88M	2805	73.6 (+4.1)	80.3 (+5.1)	60.0 (+4.7)	24.0 (+4.1)	50.7 (+6.7)
ViT-B/16-SAM	87M	863	79.9 (+5.3)	85.2 (+5.4)	67.5 (+6.2)	26.4 (+6.3)	56.5 (+9.9)
MLP-Mixer							
Mixer-S/32-SAM	19M	11401	66.7 (+2.8)	73.8 (+3.5)	52.4 (+2.9)	18.6 (+2.7)	39.3 (+4.1)
Mixer-S/16-SAM	18M	4005	72.9 (+4.1)	79.8 (+4.7)	58.9 (+4.1)	20.1 (+4.2)	42.0 (+6.4)
Mixer-S/8-SAM	20M	1498	75.9 (+5.7)	82.5 (+6.3)	62.3 (+6.2)	20.5 (+5.1)	42.4 (+7.8)
Mixer-B/32-SAM	60M	4209	72.4 (+9.9)	79.0 (+10.9)	58.0 (+10.4)	22.8 (+8.2)	46.2 (12.4)
Mixer-B/16-SAM	59M	1390	77.4 (+11.0)	83.5 (+11.4)	63.9 (+13.1)	24.7 (+10.2)	48.8 (+15.0)
Mixer-B/8-SAM	64M	466	79.0 (+10.4)	84.4 (+10.1)	65.5 (+11.6)	23.5 (+9.2)	48.9 (+16.9)

Vision Transformer (ViT)

ViT v.s. ResNet

- Let's compare one ViT layer vs one convolution layer
- Reception field: (which input neurons can affect an output neuron)
 - CNN: some subarea of image (kernel size)
 - Self-attention: the whole image
 - \Rightarrow there exists self-attention function that cannot be captured by convolution
- Is the function set of self-attention strictly larger than convolution?
 - Yes, given enough attention heads

Vision Transformer (ViT)

How can self-attention do convolution?

- Consider self-attention with relative positional encoding

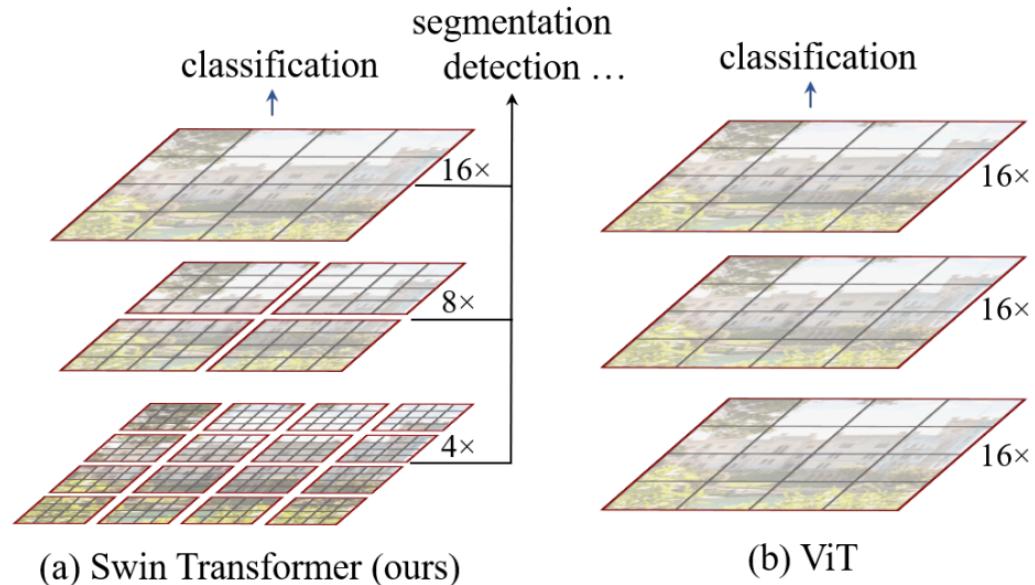
$$\text{Output} = \text{Softmax}\left(\underbrace{\frac{QK^T}{\sqrt{d}}}_{\begin{array}{c} \text{context aware} \\ \cdot \end{array}} + \underbrace{B}_{\begin{array}{c} \text{context agnostic} \\ \cdot \end{array}}\right)V$$

- Q, K, V : query, key, value matrices
- $B_{i,j} = b_{(x_i-x_j, y_i-y_j)}$: relative positional encoding (trainable scalars)
- To perform convolution: Set $Q, K = 0$ and purely rely on B
- Implication: the positional encoding can capture CNN; the query/key matrices can capture context-aware information beyond convolution

Vision Transformer (ViT)

Swin Transformer (Liu et al., 2021)

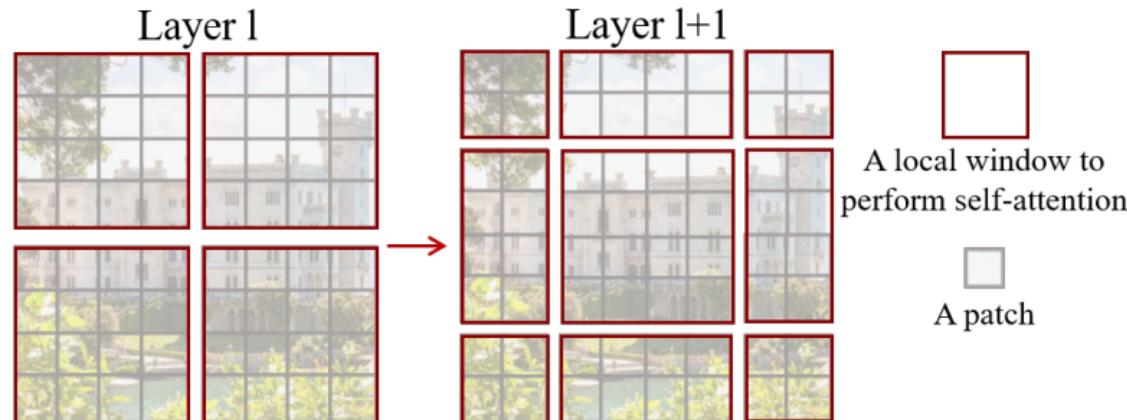
- Problems of the original ViT:
 - Non-overlapping partition
 - Only a single resolution
 - Quadratic complexity for attention computation
- Swin Transformer: hierarchical and sliding window partitions



Vision Transformer (ViT)

Swin Transformer

- Attention within sub-blocks with shifts to avoid huge attention matrix



Vision Transformer (ViT)

Swin Transformer

(a) Regular ImageNet-1K trained models					
method	image size	#param.	FLOPs	throughput (image / s)	ImageNet top-1 acc.
RegNetY-4G [48]	224 ²	21M	4.0G	1156.7	80.0
RegNetY-8G [48]	224 ²	39M	8.0G	591.6	81.7
RegNetY-16G [48]	224 ²	84M	16.0G	334.7	82.9
EffNet-B3 [58]	300 ²	12M	1.8G	732.1	81.6
EffNet-B4 [58]	380 ²	19M	4.2G	349.4	82.9
EffNet-B5 [58]	456 ²	30M	9.9G	169.1	83.6
EffNet-B6 [58]	528 ²	43M	19.0G	96.9	84.0
EffNet-B7 [58]	600 ²	66M	37.0G	55.1	84.3
ViT-B/16 [20]	384 ²	86M	55.4G	85.9	77.9
ViT-L/16 [20]	384 ²	307M	190.7G	27.3	76.5
DeiT-S [63]	224 ²	22M	4.6G	940.4	79.8
DeiT-B [63]	224 ²	86M	17.5G	292.3	81.8
DeiT-B [63]	384 ²	86M	55.4G	85.9	83.1
Swin-T	224 ²	29M	4.5G	755.2	81.3
Swin-S	224 ²	50M	8.7G	436.9	83.0
Swin-B	224 ²	88M	15.4G	278.1	83.5
Swin-B	384 ²	88M	47.0G	84.7	84.5

(b) ImageNet-22K pre-trained models					
method	image size	#param.	FLOPs	throughput (image / s)	ImageNet top-1 acc.
R-101x3 [38]	384 ²	388M	204.6G	-	84.4
R-152x4 [38]	480 ²	937M	840.5G	-	85.4
ViT-B/16 [20]	384 ²	86M	55.4G	85.9	84.0
ViT-L/16 [20]	384 ²	307M	190.7G	27.3	85.2
Swin-B	224 ²	88M	15.4G	278.1	85.2
Swin-B	384 ²	88M	47.0G	84.7	86.4
Swin-L	384 ²	197M	103.9G	42.1	87.3